



# Recommender Systems

Noopur Khachane  
Mihir Sathe  
Liyang Yu  
Heguang Zhou  
Jingpei Lu



# What is a Recommender System?

- Given data about users and items, predict a user's preferences or ratings for items

# Where do we see them today?


- Recommender systems are so common nowadays that we probably don't even think about them.
  - Amazon's "Suggested Items"
  - Spotify "Discover Playlist"
  - Netflix's "Match Score"

# Netflix Prize

- From the years 2006-2009, Netflix offered up a prize of 1 million dollars to a team that could beat their recommender system by 10%
- The winners used combinations of many algorithms, one team even used about 107.
- The challenge was to predict without information about users

# MovieLens100K

- Website that supplies movie recommendations
- Open-Source dataset
- Includes (100,000)
  - User Info(Age, Gender, ZIP, Occupation)
  - Movie Info(Release date, Genre)



# Collaborative Filtering

The background of the slide features a light gray world map. Overlaid on the map is a complex network of nodes and lines. The nodes are represented by circles of varying sizes and shades of gray, some with concentric circles. The lines are thin and gray, connecting the nodes across the map, suggesting a global network or data flow. The overall aesthetic is technical and modern.

- Find 'Neighbors' with similar preferences to target user
- Use a similarity weighted average to predict
- Eg. K-Nearest Neighbors

# Content-Based Filtering

- Weight all movies with respect to similarity with the active movie.
- Select a subset of the movies(*neighbors*) to use as predictors.
- Compute a prediction from a weighted combination of the selected neighbors' ratings.


# Hybrid Systems

- Take information from Collaborative/Content Filtering
- Compute a prediction from a weighted combination of the two models



# Neural Network

- Engineer Features based on (Users, Movies, Ratings)
- Train Network on (Features, Ratings)
- Predict for new (User,Movie)



# Process Data First!

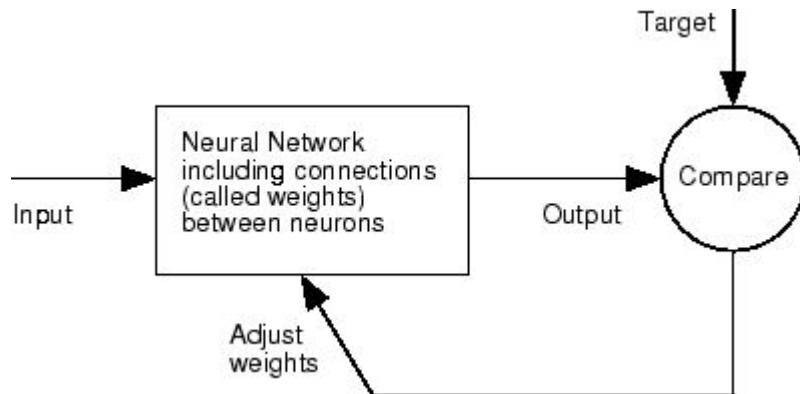
- Split into Training/Test with 80:20 Ratio
- 80,000 Training Examples, 20,000 Test Examples

# Neural Network

- Engineer Features based on (Users, Movies, Ratings)
  - `three_year_mean(Users,Movies,Ratings)`
  - `gender_mean(Users,Movies,Ratings)`
  - `user_mean(User,Ratings)`
  - `movie_mean(Movie,Ratings)`

# Neural Network

- Train on 80,000 (Features, Ratings)
  - The network learns to predict a rating based on a function of its features



# Neural Network

- Predict on 20,000 Test Points
- Calculate Root Mean Squared Error = 0.85★
  - “How many stars are we ‘off’ by?”
- Calculate Explained Variance = 67.5%
  - “What percent of the variation in our dataset does our model explain?”

# Future Improvements

- Predict on more data (MovieLens1M)
- Extract more features
  - We threw away ZIP, Occupation
- Optimize Model Parameters
  - How to choose size of network? Learning Rate? etc.

Thank You!